

Search and Motion Planning

Prof. Dr. Daniel Göhring

January 29, 2018

FOR INTERNAL USE ONLY

Table of contents

1 Planning Basics

- State Space Search, A^*
- Heuristic and Costs
- Alternative Forms of A^*

State Space Search¹

- States Z (Vertices)
- Operators $Op \subseteq Z \times Z$ (Edges)
- Graph $[Z, Op]$
- Initial states $z_{initial} \in Z$
- Target states $Z_{final} \subseteq Z$
- Cost functions $c : Z \times Z \rightarrow R^+$,
 Costs of path $w = z_0 z_1 \cdots z_n \in Z^*$:
 $c(z_0 z_1 \cdots z_n) := \sum_{i=1, \dots, n} c(z_{i-1}, z_i)$
- Estimator function $\sigma : Z \rightarrow R$ (Heuristic) for remaining from z to a target state.

¹Slides with courtesy of Prof. Dr. Hans-Dieter Burkhard [2]

State Space Search

- Tasks:
- Can a goal state $z \in Z_{final}$ be reached from initial state $z_{initial} \in Z$?
- Find a way from initial state $z_{initial} \in Z$ to a target state $z \in Z_{final}$
- Find an optimal path from initial state $z_{initial} \in Z$ to a target state $z \in Z_{final}$

Complexity (Number of States / Vertices)


- 8 Puzzle: $9!$ states
 $9!/2 = 181.440$ reachable
- 15 Puzzle: : $16!$ states
 $16!/2$ reachable
- Rubik's cube: $12 \cdot 4.3 \cdot 10^{19}$ states
 $1/12$ reachable: $4.3 \cdot 10^{19}$
- Towers of Hanoi: 3^n States for n discs
solvable in $(2^n) - 1$ moves
- Checkers: approx. 10^{40} games of average length
- Chess: approx. 10^{120} games of average length
- Go: 3^{361} states

Expansion Strategies

- Directions
 - Forward, start with $z_{initial}$ (forward chaining, data driven, bottom up)
 - Reverse z_{final} (backward chaining, goal driven, top down)
 - Bi-directional
- Expansion
 - Depth first
 - Breadth first
- Additional information
 - blind search ("uninformed")
 - heuristic search with σ ("informed")

Bi-Directional Breadth-Search

- Parallel search from start and goal until meeting



- Search depth from both sides only half

Expansion

Data structures:

- OPEN List: A vertex is "open" , if it was constructed but not expanded (neighboring vertices not calculated)
- CLOSED List: A vertex is "closed" , if it was fully expanded (all neighboring vertices are known)
- Further information: Predecessor / successor of vertices for reconstruction of found paths

Heuristic Search for best Way: A*

Costs to reach z' from z :

- If z' is reachable from z :

$$g(z, z') := \min\{c(s) \mid s \text{ path from } z \text{ to } z'\}$$


- Else: $g(z, z') := \infty$

Tentative cost calculation during expansion: $G' = [V', E']$ is a (known) partial graph of G

$$g'(z, z', G') := \min\{c(s) \mid s \text{ path in } G' \text{ from } z \text{ to } z'\}$$

$$g'(z, z', G') \geq g(z, z')$$


Heuristic Search for Best Path



$g'(z_0, z', G')$: so far known costs to reach z' from Start

$\sigma(z')$: estimated costs to reach target state, starting from z'

State Space Search - Example



Algorithm A* (soft form) for Trees

- A*0 (Start) $OPEN := [z_0]$, $CLOSED := []$.
- A*1: (negative exit)
If $OPEN = []$: EXIT("no").
- A*2: (positive exit)
If z first vertex in $OPEN$:
If z is target: EXIT("yes:" z).
- A*3: (expand)
 $OPEN := OPEN - z$. $CLOSED := CLOSED \cup \{z\}$
 $Succ(z) :=$ set of successors of z .
If $Succ(z) = \{\}$: Goto A*1
- A*4: (Organization of $OPEN$)
 $OPEN := OPEN \cup Succ(z)$ with increasing order of
 $g'(z_0, z', G') + \sigma(z')$
- Goto A*1. Search space as a tree: CLOSED not used.

Algorithm A* (soft form) for Trees and for Cyclic Graphs

- A*0 (Start) $OPEN := [z_0]$, $CLOSED := []$.
- A*1: (negative exit)
If $OPEN = []$: EXIT("no").
- A*2: (positive exit)
If z first vertice in $OPEN$:
If z is target: EXIT("yes:" z).
- A*3: (expand)
 $OPEN := OPEN - z$. $CLOSED := CLOSED \cup \{z\}$
 $Succ(z) :=$ set of successors of z .
If $Succ(z) = \{\}$: Goto A*1
 $NEW = Succ(z) - \{z' | z' \in Succ(z) \text{ and } z' \in CLOSED \text{ and } g'(z_0, z', G') \geq g'(z_0, z', G'_{old})\}$

Algorithm A* (soft form) for Trees and for Cyclic Graphs

- A*4: (Organization of *OPEN*)
 $OPEN := OPEN \cup NEW$ with increasing order of
 $g'(z_0, z', G') + \sigma(z')$
- Goto A*1. If search space has cycles: use CLOSED.

Algorithm A* ("soft form")

Definition

$f(z) := \min\{g(z, z_{final}) \mid z_{final} \in Z_{final}\}$ = real costs from z to target state (vertex)

$(f(z_0) = \text{cost of the optimal path})$

Heuristic function σ is called optimistic (a.k.a. admissible) or underestimating, if $\sigma(z) \leq f(z)$ for all $z \in Z$.

Proposition

Given: Ex. $\delta > 0$ with $c(z, z') > \delta$ for all z, z' . σ is an optimistic / admissible heuristic. Every node has a finite number of successors. One can prove: If solution exists, A* (soft form) finds an optimal path.

Special Cases

$c \equiv 0$:

Search for best path with heuristic σ and no cost function (Hill climbing)

$\sigma \equiv 0$:

Search for best path without heuristic ($s \equiv 0$ is also optimistic heuristic) - Dijkstra

$c \equiv 1$

($g' \equiv$ Search depth) , $\sigma \equiv 0$: Breadth first search

Influence of Heuristic σ


- Trade-off between quality of result and calculation effort
- Criteria for order of extension
- Same order as $g' + \sigma$ is also provided by $a(g' + \sigma) + b$ for arbitrary positive constants a, b .
- optimal order for $\sigma = f$
- σ_2 more efficient than σ_1 if $\sigma_1 \leq \sigma_2 \leq f$ (Hierarchy for heuristic functions)

Algorithm A* (hard form)

- A*0 (Start) $OPEN := [z_0]$, $CLOSED := []$.
- A*1: (negative exit)
If $OPEN = []$: EXIT("no").
- A*2: (positive exit)
If z first vertex in $OPEN$:
If z is target: EXIT("yes:" z).
- A*3: (expand)
 $OPEN := OPEN - z$. $CLOSED := CLOSED \cup \{z\}$
 $Succ(z) :=$ set of successors of z .
If $Succ(z) = \{\}$: Goto A*1
- A*4: (Organization of $OPEN$)
 $OPEN := OPEN \cup (Succ(z) - CLOSED)$ sorted with
increasing order of $g'(z_0, z', G') + \sigma(z')$
- Goto A*1.

Algorithm A* (hard form)

- Problem:
- For optimistic σ the hard form is not always correct



Algorithm A* (hard form)

What kind of heuristic do we need to stay optimal?

Definition

A heuristic σ is called consistent or monotonous, if for all states z', z'' holds: $\sigma(z') \leq g(z', z'') + \sigma(z'')$

Lemma

If σ is consistent, σ is optimistic. The reverse does not necessarily apply

Algorithm A* (hard form)

Proposition

Given: Ex. $\delta > 0$ with $c(z, z') > \delta$ for all z, z' . σ is a consistent heuristic. One can prove: if a solution exists, A* (hard form) finds an optimal path.

Algorithm A* ("hard form")

- "consistent" functions harder to find than "optimistic"
- It is also possible to use an optimistic heuristic with weaker pruning of search space
- Erase states in OPEN (or Succ(z)) only if new evaluation is worse than earlier evaluation.

Memory Saving Variants of Algorithm A*


- Iterative Deepening A* (IDA*) in analogy to IDA: Depth first till boundary $g'(z_0, z', G') + s(z')$ is exceeded from earlier iteration
- SMA* (simplified memory-bounded A*)

Anytime (or Weighted) A*

- If the heuristic σ (sometimes denoted as h) is closer to the real costs, less vertices have to be expanded
- But "inflating" the heuristics can lead to a heuristic which is not optimistic (admissible)

Heuristics in Heuristic Search²

- Dijkstra's: expands states in the order of $f = g$ values
- A* Search: expands states in the order of $f = g + \sigma$ values
- Weighted A*: expands states in the order of $f = g + \varepsilon\sigma$ values, $\varepsilon > 1$ = bias towards states that are closer to goal




²Slides from Maxim Likhachev

Suboptimality

- Heuristic includes factor ε
- Suboptimality is bounded by factor ε
 - The length of the found solution is not longer than $\frac{1}{1-\varepsilon}$ times the optimal solution
- Exemple:
 - costs from cell to cell are 1
 - Heuristic is the larger of coordinate difference from current cell to goal cell
 - Start is upper left cell
 - Goal is lower right
 - Obstacles black
 - Free space white
 - Expanded cells gray

Anytime Search based on weighted A*

- Constructing anytime search based on weighted A*:
 - find the best path possible given some amount of time for planning
 - do it by running a series of weighted A* searches with decreasing ϵ :



Anytime A*

- Problem:
 - Running A* with increasing ε each time from scratch can be very expensive
 - Many states remain the same through various iterations
 - A solution for reusing the search results is described in ARA* (Likhachev [4])
- ARA* : an efficient version of the above that reuses state values within any search iteration

ARA*

- Efficient series of weighted A* searches with decreasing ε :

ComputePathwithReuse function

```

while( $f(s_{goal}) > \text{minimum } f\text{-value in } OPEN$ )
  remove  $s$  with the smallest  $[g(s) + \varepsilon h(s)]$  from  $OPEN$ ;
  insert  $s$  into  $CLOSED$ ;
  for every successor  $s'$  of  $s$ 
    if  $g(s') > g(s) + c(s, s')$ 
       $g(s') = g(s) + c(s, s')$ ;
    if  $s'$  not in  $CLOSED$  then insert  $s'$  into  $OPEN$ ;
    otherwise insert  $s'$  into  $INCONS$ 

```


```

set  $\varepsilon$  to large value;
 $g(s_{start}) = 0$ ;  $OPEN = \{s_{start}\}$ ;
while  $\varepsilon \geq 1$ 
   $CLOSED = \{\}$ ;  $INCONS = \{\}$ ;
  ComputePathwithReuse();
  publish current  $\varepsilon$  suboptimal solution;
  decrease  $\varepsilon$ ;
  initialize  $OPEN = OPEN \cup INCONS$ ;


```

ARA*, slides from Likhachev [4]

- A series of weighted A* searches




- ARA*



D* and Variations

- D* and its variations have been used on Mars rovers Opportunity and Spirit and by CMU at the DARPA Grand Challenge
- Roughly:
 - D* works as A* but from goal to start
 - Every expanded node "knows" its predecessor
 - When start node (vertex) is the next node, the search is done
 - nodes are marked: NEW (was never in OPEN), OPEN, CLOSE (no longer in OPEN), LOWER, RAISE (cost is higher than last time in OPEN)

D* and Variations




Heuristics in Heuristic Search [4]

Heuristics in Heuristic Search

- Dijkstra's: expands states in the order of $f = g$ values

What are the states expanded?




Heuristics in Heuristic Search [4]

Heuristics in Heuristic Search

- A* Search: expands states in the order of $f = g + h$ values


for high-D problems, this results in A being too slow and running out of memory*



Heuristics in Heuristic Search [4]

Heuristics in Heuristic Search

- Weighted A* Search: expands states in the order of $f = g + \varepsilon h$ values, $\varepsilon > 1$
- bias towards states that are closer to goal



Literature

-  [1] J.C. Latombe, Robot Motion Planning, Kluwer Academic Publishers, 1991.
-  [2] Hans-Dieter Burkhard: Einführung in die Künstliche Intelligenz (2008)
-  [3] S. LaValle, Planning Algorithms. 2006.
<http://msl.cs.uiuc.edu/planning/>
-  [4] Likhachev, ARA*, CMU
-  [5] Choset, Motion Planning, CMU,
-  [6] Toussaint, Lecture Notes Robotics, 2011
-  [7] Dr. John (Jizhong) Xiao, City College New York